

Section 16: Planning and configuring storage

Storage overview

- object storage
 - ↳ cloud storage
- Instance storage
 - ↳ persistent disks
- Database storage
 - SQL
 - ↳ cloud SQL
 - ↳ cloud Spanner
 - No SQL
 - ↳ Cloud Data / File store → Docs
 - ↳ Bigtable → columnar
 - Analytic
 - ↳ Big Query

Cloud Storage

- unstructured Data
 - ↳ Images + video + Text
- Archived Data
- Temporary storage [moving data]

Local organization

- Global Name Space
 - ↓
 - Buckets
 - ↓
 - Folders

Classes

- Standard storage
 - * frequent access
 - * performance optimized
- Archive + Backup storage
 - ↳ Nearline
 - ↳ stored at least 30 days
 - ↳ Cold line
 - ↳ at least 90 days
 - ↳ Archive
 - ↳ at least 365 days
- * there is also a read/write guidance once per x days
 $x = 30/90/365$

Redundancy

- Regional
 - ↳ single Region
 - ↳ Replicated across zones
- Dual-Region
 - ↳ 2 Regions
 - ↳ auto failover
- Multi-Region
 - ↳ US, Europe, Asia
 - ↳ auto-failover

Object Lifecycle Policies

- auto-automatic transition to cheaper tier
- Retention Policies (time)
 - ↳ minimum period
- Hold delete prevention

Security

- All encrypted
 - google keys
 - customer cloud key
 - customer supplied
- Bucket access can be managed with
 - ↳ Bucket level Rules
 - ↳ IAM
- auto-atic logging

Cloud SQL

- Hardware options
- auto backups
- High availability
 - ↳ replication
- MySQL Database Plugins
- Maintenance Schedule

Database Info can be found in instance details

Cloud Spanner

- Scalable Horizontally
- Strong consistency
- availability around the world

Together with SQL Relational db options

Pick Multi-Zone deployment
Spanner is expensive!

Cloud Data Store

- NoSQL Database
- Modes
 - Native Mode
 - DataStore Mode

Serverless Service

One Database per project

- DataStore is created automatically
 - ↳ Api Based
 - ↳ entities cost Money

Entity

properties can be different per Row because of NoSQL

Cloud Firestore

- Managed NoSQL
- Document Based
 - ↳ key-value
 - ↳ Hierarchical
 - ↳ json like

Use when

- ↳ Semi-structured Data
 - ↳ Not Fixed Schema
 - ↳ Different attributes across schema

- Query on Multiple attribute
- Ingestion volume not too high

2 Modes

- DataStore Mode
 - * Backend for server apps
 - * No Syncing with Mobile device
- Native Mode
 - * Mobile + web apps
 - * Large Number of connections

Cloud Bigtable

- No SQL
- Wide-column
- Low-latency + high availability
- works with analytics
- Not Serverless

Bigtable is expensive!

Type

- Production
- Development

Hardware options

Node based

- ↳ Metadata is stored in nodes telling services where data is located

For:

- * High volume
- * High velocity

Based on colossus

Cloud MemoryStore

- ↳ Caching Data
 - ↳ DB's
 - ↳ Games
 - ↳ Data streams

→ Redis

- * Connect from lots of different services

- * up to 300 gb
- * 12 Gbps throughput

Redis Protocol Compliance Service

- Basic Tier
 - no replication
 - no failover
- Standard tier
 - Replication
 - failover

→ MemCached

In-Memory key-value store
↳ Distributed

For:

Reference data
Query cache + cache results
Session cache

- Is one cluster
 - ↳ Max 20 nodes
 - ↳ Max 5 Tb memory
 - ↳ so good for large cache needs

→ supports different services.